



电子科技大学
格拉斯哥学院
Glasgow College, UESTC

UESTC4019: Real-Time Computer Systems and Architecture

Lecture 17

Instruction-Level Parallelism, and Superscalar Processors (Part-1)

Superscalar Overview


- Term first coined in 1987
- **Superscalar** refers to a machine that is designed to improve the performance of the execution of scalar instructions
- In most applications the bulk of the operations are on scalar quantities
- Represents the next step in the evolution of high-performance general-purpose processors
- Essence of the approach is the ability to execute instructions independently and concurrently in different pipelines
- Concept can be further exploited by allowing instructions to be executed in an order different from the program order

Superscalar Organization Compared to Ordinary Scalar Organization


In a **traditional scalar organization**, there is a **single pipelined functional unit** for integer operations and one for floating-point operations

Parallelism is achieved by enabling multiple instructions to be **at different stages of the pipeline at one time**

In the **superscalar organization**, there are **multiple functional units**, each of which is implemented as a pipeline



(a) Scalar organization



(b) Superscalar organization

Superscalar vs Ordinary Scalar

- Each individual functional unit provides a degree of **parallelism by virtue of its pipelined structure**.
- The use of multiple functional units enables the processor **to execute streams of instructions in parallel, one stream for each pipeline**.
- It is the responsibility of the hardware, in conjunction with the compiler, to assure that the **parallel execution does not violate the intent of the program**.

Superpipelining

- An alternative approach to achieving greater performance is referred to as super-pipelining, a term first coined in 1988
- **Superpipelining** exploits the fact that many pipeline stages perform tasks that require less than half a clock cycle
- Thus, a doubled internal clock speed allows the performance of two tasks in one external clock cycle
- We have seen one example of this approach with the MIPS R4000


Comparison of Superscalar and Super Pipeline Approaches (1 of 3)

The upper part of the diagram illustrates an **ordinary pipeline**, used as a base for comparison

The **base pipeline** issues one instruction per clock cycle and can perform one pipeline stage per clock cycle

The **pipeline has four stages**: instruction **fetch**, operation **decode**, operation **execution**, and result **write** back. The execution stage is crosshatched for clarity.

Note that although several instructions are executing **concurrently**, only one instruction is in its execution stage at any one time.

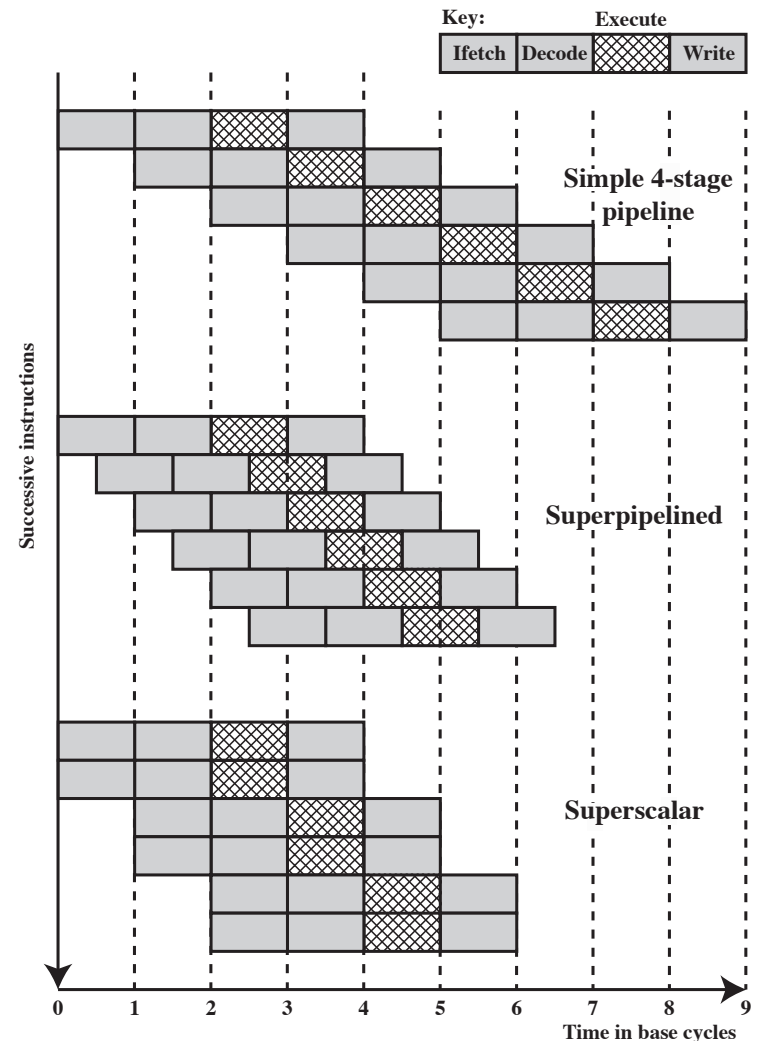


Comparison of Superscalar and Super Pipeline Approaches (2 of 3)

The next part of the diagram shows a **superpipelined** implementation that is **capable of performing two pipeline stages per clock cycle**.

An alternative way of looking at this is that the functions performed in each stage **can be split into two non-overlapping parts** and **each can execute in half a clock cycle**.

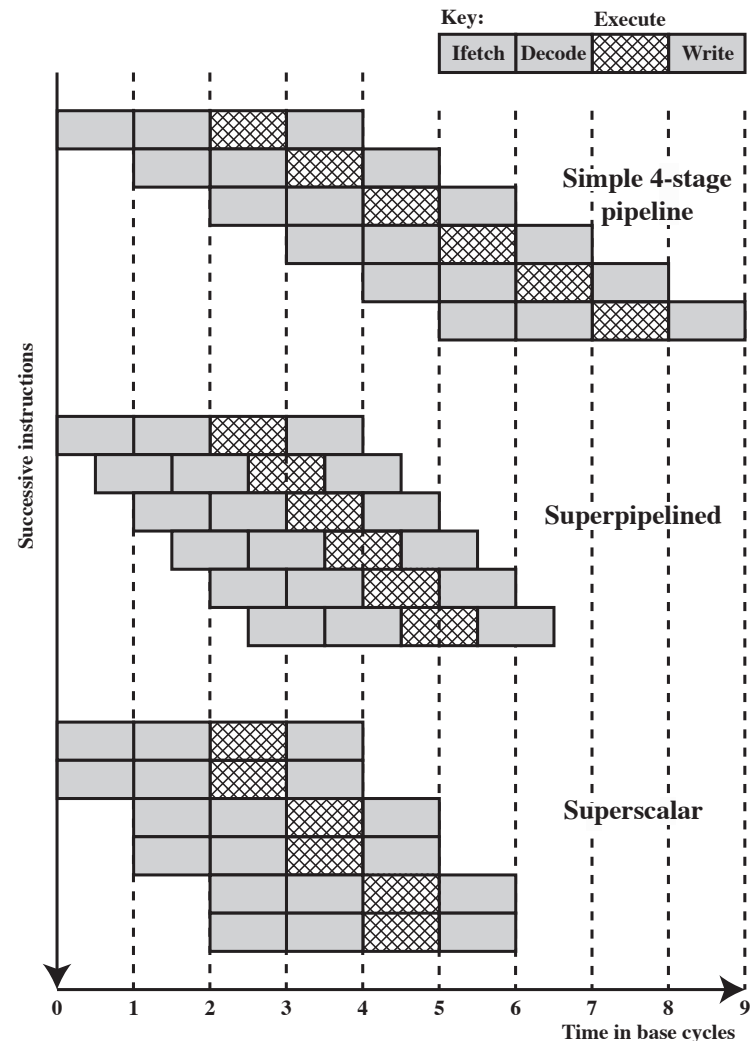
A superpipeline implementation that **behaves** in this fashion is said to be of **degree 2**.



Comparison of Superscalar and Super Pipeline Approaches (3 of 3)

Finally, the lowest part of the diagram shows a **superscalar implementation** capable of **executing two instances of each stage in parallel**.

Higher-degree superpipeline and super-scalar implementations are of course possible.



Constraints

- Instruction level parallelism
 - Refers to the degree to which the instructions of a program can be executed in parallel
 - A combination of compiler based optimization and hardware techniques can be used to maximize instruction level parallelism
- Limitations
 - True data dependency
 - Procedural dependency
 - Resource conflicts
 - Output dependency
 - Antidependency

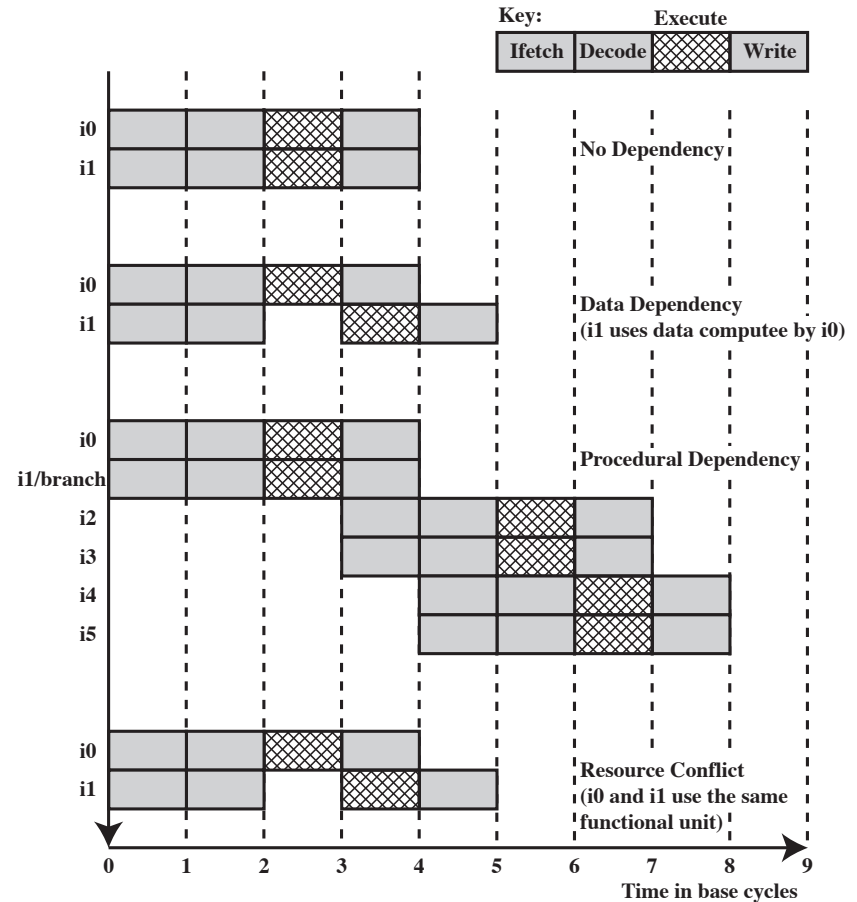
Effect of Dependencies

The presence of **branches** in an instruction sequence complicates the pipeline operation

The instructions following a branch (taken or not taken) have a **procedural dependency** on the branch and **cannot be executed until the branch is executed**

Figure illustrates the effect of a branch on a superscalar pipeline of degree 2

The **consequence of procedural dependency** for a superscalar pipeline is more severe, because a greater magnitude of opportunity is lost with each delay



Design Issues

- Instruction level parallelism
 - Instructions in a sequence are independent
 - Execution can be in parallel by overlapping
 - Governed by data and procedural dependency
- Machine Parallelism
 - Ability of the processor to take advantage of instruction level parallelism
 - Governed by number of parallel pipelines

Instruction Issue Policy (1 of 2)

- Instruction issue
 - Refers to the process of **initiating instruction execution in the processor's functional units**
- Instruction issue policy
 - Refers to the **protocol used to issue instructions**
 - Instruction issue occurs when instruction moves from **the decode stage** of the pipeline to **the first execute stage** of the pipeline

Instruction Issue Policy (2 of 2)

- Three types of orderings are important:
 - The order in which instructions are fetched
 - The order in which instructions are executed
 - The order in which instructions update the contents of register and memory locations
- Superscalar instruction issue policies can be grouped into the following categories:
 - In-order issue with in-order completion
 - In-order issue with out-of-order completion
 - Out-of-order issue with out-of-order completion

Superscalar Instruction Issue and Completion Policies

Figure gives an example of this policy.

We assume a **superscalar pipeline** capable of fetching and decoding two instructions at a time, having three separate functional units (e.g., two integer arithmetic and one floating-point arithmetic), and having two instances of the write-back pipeline stage.

The example assumes the following constraints on a six-instruction code fragment:

Decode		Execute			Write		Cycle
I1	I2						1
I3	I4	I1	I2				2
I3	I4	I1					3
	I4			I3	I1	I2	4
I5	I6			I4			5
	I6		I5		I3	I4	6
			I6				7
					I5	I6	8

(a) In-order issue and in-order completion

Decode		Execute			Write		Cycle
I1	I2						1
I3	I4	I1	I2				2
	I4	I1		I3	I2		3
I5	I6			I4	I1	I3	4
	I6		I5		I4		5
			I6		I5		6
					I6		7

(b) In-order issue and out-of-order completion

Decode		Window	Execute			Write		Cycle
I1	I2	I1,I2	I1	I2				1
I3	I4	I3,I4	I1		I3	I2		2
I5	I6	I4,I5,I6		I6	I4	I1	I3	3
		I5		I5		I4	I6	4
						I5		5
								6

(c) Out-of-order issue and out-of-order completion

Register Renaming

- Output and anti dependencies occur because **register contents** may not reflect **the correct ordering** from the program
- May result in a pipeline stall
- Registers allocated dynamically

Branch Prediction (1 of 2)

- Any high-performance **pipelined machine** must address the issue of **dealing with branches**
- Intel 80486 addressed the problem by **fetching both** the **next sequential** instruction after a branch and **speculatively fetching** the branch target instruction
- RISC machines:
 - **Delayed branch** strategy was explored
 - Processor always **executes the single instruction** that immediately follows the branch

Branch Prediction (2 of 2)

- Keeps the **pipeline full** while the processor fetches a new instruction stream
- Superscalar machines:
 - **Delayed branch** strategy has less appeal
 - Have returned to pre-RISC techniques of branch prediction